

# **How to Handle 10,000 Users in 1 Second**

**Solving the “Taylor Swift Problem”  
with Node.js & Redis**

Kavishka Niluminda

---

## The Logic



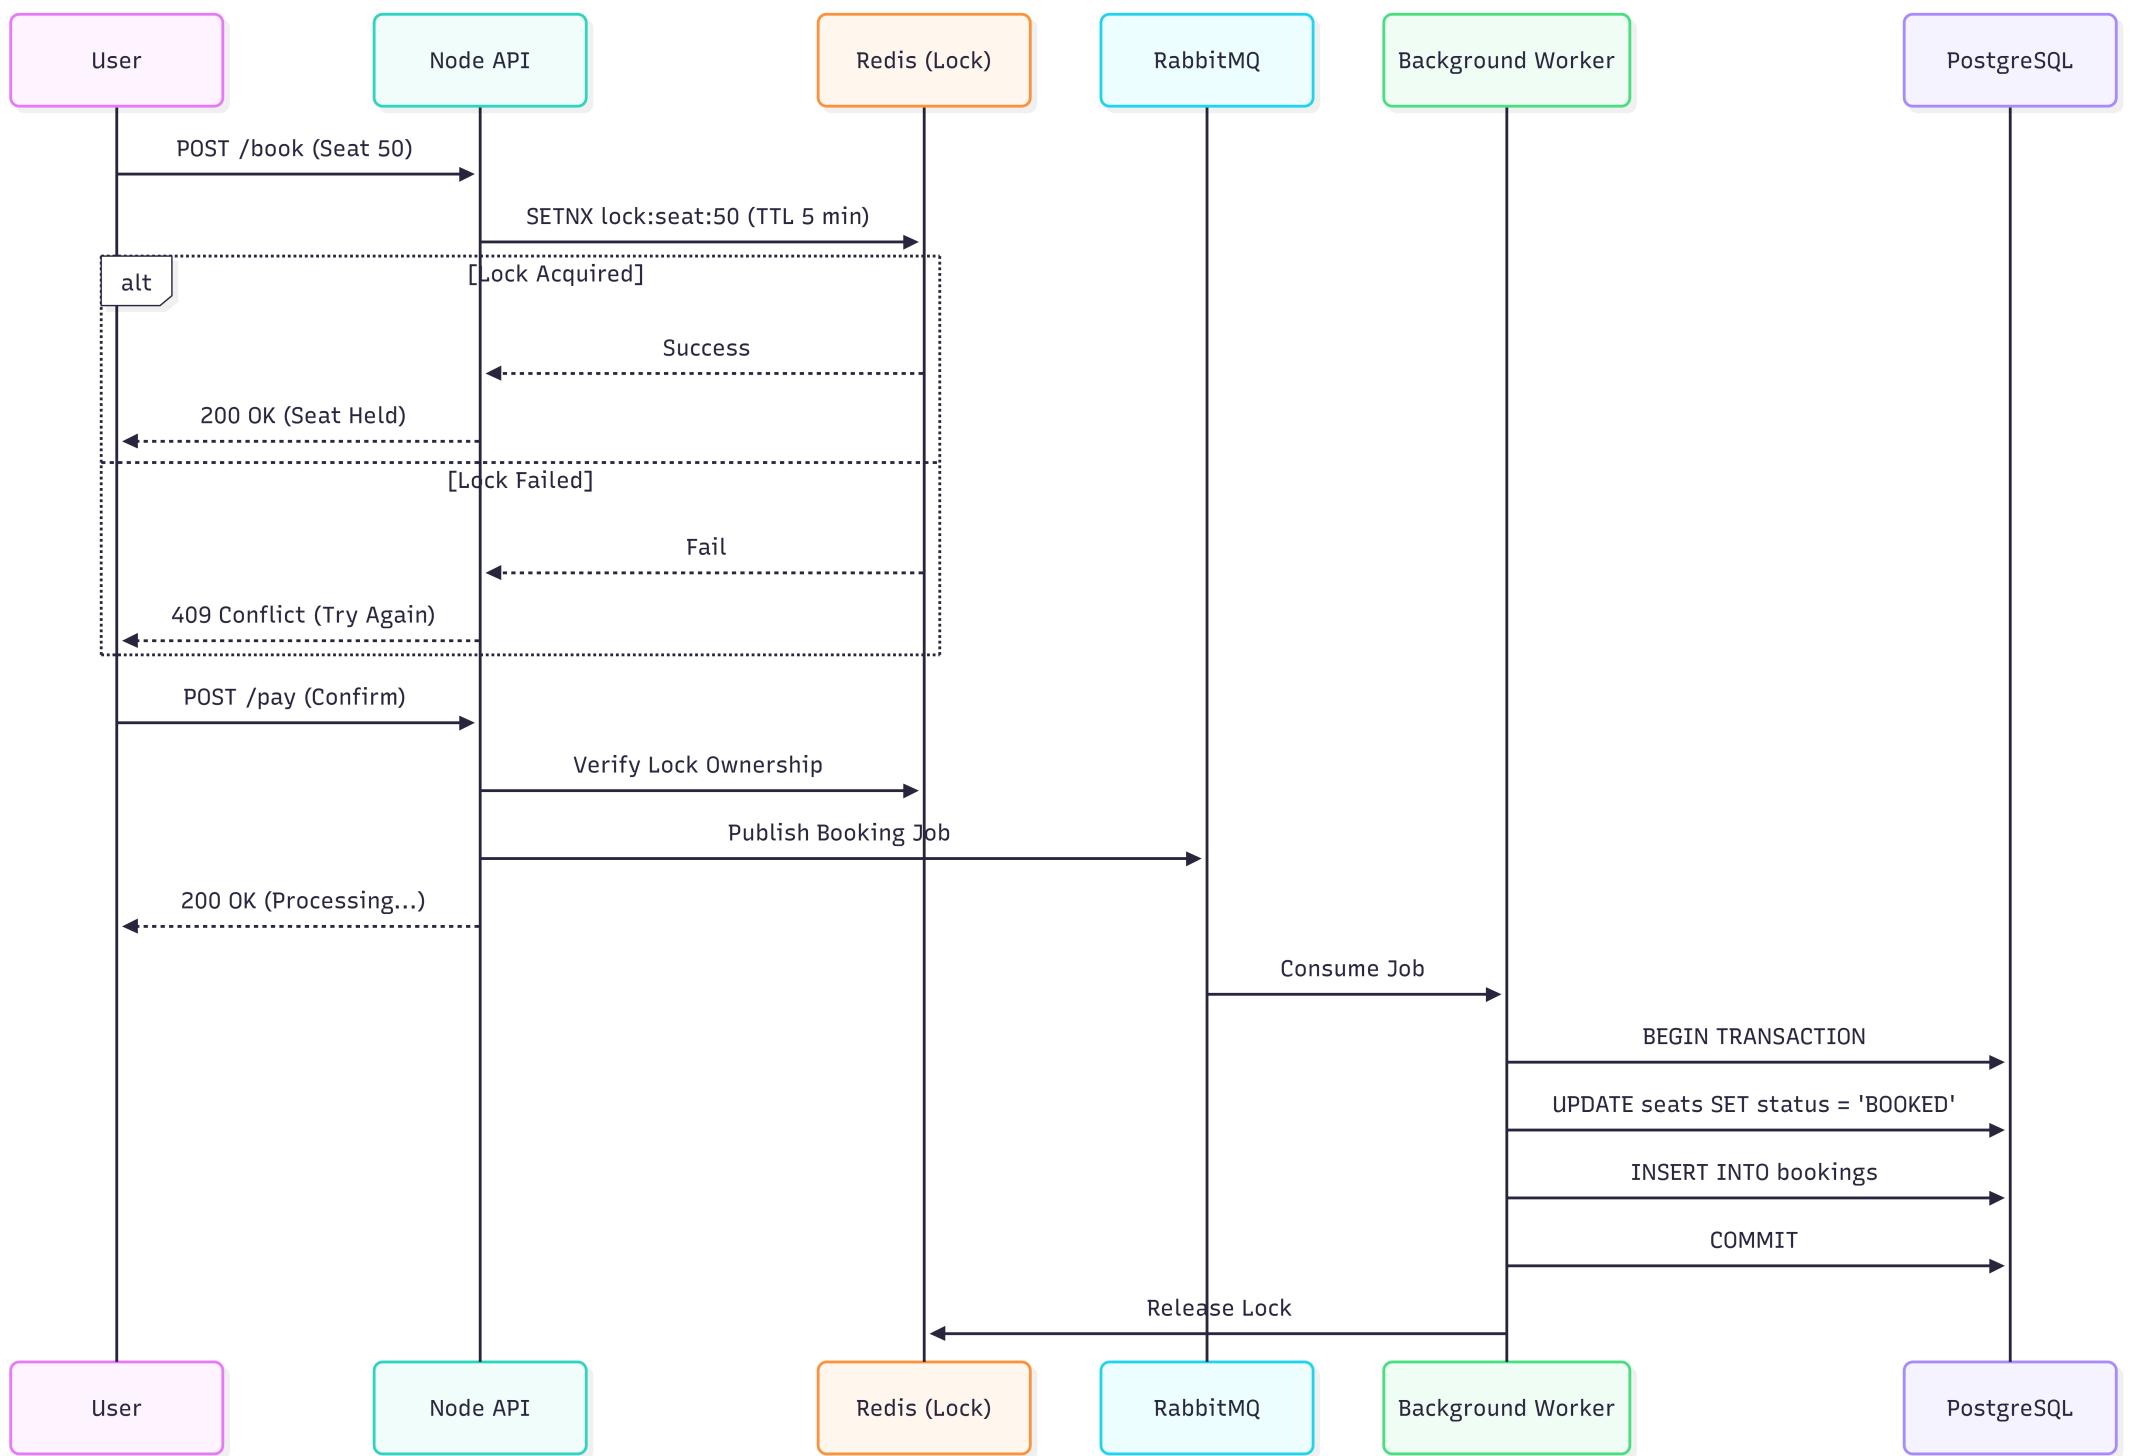
```
// THE CORE LOGIC: Distributed Locking with Redis
// Goal: Prevent double-booking when 10k users hit the button at once.

export const bookSeat = async (req: Request, res: Response) => {
  const { seatId, userId } = req.body;
  const lockKey = `lock:seat:${seatId}`;

  try {
    // 1. Try to acquire the lock in Redis (Mutex)
    // 'NX': Only set if Not Exists (The Atomic Guard)
    // 'EX': Auto-expire in 300s to prevent deadlocks
    const acquiredLock = await redisClient.set(lockKey, userId, {
      NX: true,
      EX: 300
    });

    // 2. If Redis says "False", someone else beat us by a millisecond
    if (!acquiredLock) {
      return res.status(409).json({
        error: "Seat is currently on hold by another user."
      });
    }

    // 3. Lock Acquired! The user now has exclusive access
    // We can now safely proceed to payment processing...
    return res.status(200).json({
      message: "Seat held! Proceed to payment."
    });
  } catch (error) {
    console.error("Concurrency Error:", error);
    res.status(500).json({ error: "System Busy" });
  }
};
```



# K6 Results



## TOTAL RESULTS

```
checks_total.....: 1000  97.186916/s
checks_succeeded.: 50.00% 500 out of 1000
checks_failed....: 50.00% 500 out of 1000

X is status 200 (Success)
↳ 0% - ✓ 1 / X 499
X is status 409 (Locked)
↳ 99% - ✓ 499 / X 1
↳ 99% - ✓ 499 / X 1

HTTP
http_req_duration.....: avg=20.76ms min=551.2µs med=13.2ms max=75.1ms p(90)=57.86ms
p(95)=63.47ms
  { expected_response:true }....: avg=39.83ms min=39.83ms med=39.83ms max=39.83ms p(90)=39.83ms
p(95)=39.83ms
  http_req_failed.....: 99.80% 499 out of 500
  http_reqs.....: 500    48.593458/s

↳ 99% - ✓ 499 / X 1

HTTP
http_req_duration.....: avg=20.76ms min=551.2µs med=13.2ms max=75.1ms p(90)=57.86ms
p(95)=63.47ms
  { expected_response:true }....: avg=39.83ms min=39.83ms med=39.83ms max=39.83ms p(90)=39.83ms
p(95)=39.83ms
  http_req_failed.....: 99.80% 499 out of 500
  http_reqs.....: 500    48.593458/s

EXECUTION
iteration_duration.....: avg=1.02s   min=1s      med=1.01s   max=1.08s   p(90)=1.06s
p(95)=1.07s
  iterations.....: 500    48.593458/s
  vus.....: 50      min=50      max=50
  vus_max.....: 50      min=50      max=50

NETWORK
data_received.....: 160 kB 16 kB/s
data_sent.....: 79 kB 7.7 kB/s

running (10.3s), 00/50 VUs, 500 complete and 0 interrupted iterations
default ✓ [=====] 50 VUs 10s
```